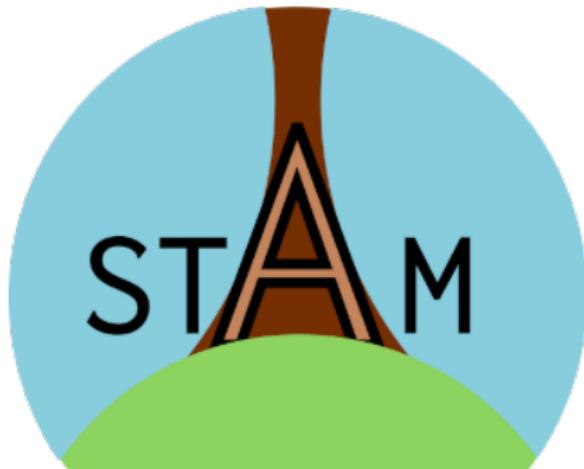


STAM: Working with stand-off annotations on text

Maarten van Gompel

KNAW Humanities Cluster

December 2025



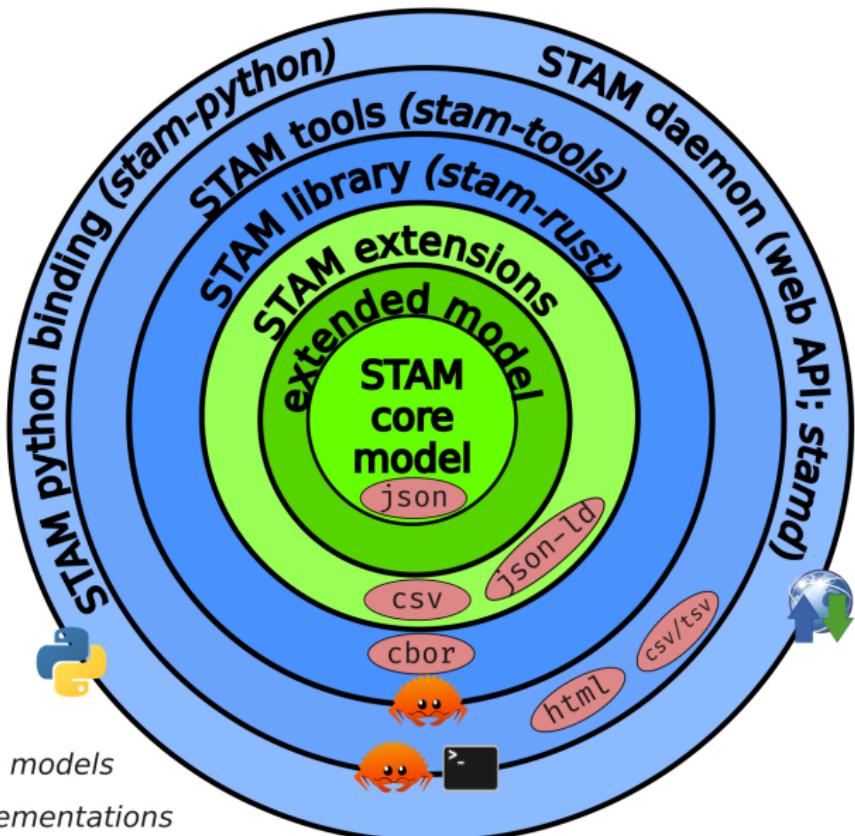
Introduction

What is STAM?

1. A **data model** for **stand-off annotation** on **text**
2. A set of *low-level* **software tools** for working with **stand-off annotation** on **text**

*Theory & Practise go hand in hand, STAM aims to provide a *solid generic foundation upon which you can build.**

Project website: <https://annotation.github.io/stam>



Data models



Implementations



Serialisation formats



Interfaces

Core Data Model

What basic concepts and relations are needed to express stand-off annotation on text?

- ▶ ... whilst keeping the concepts to a basic yet flexible enough **minimum** (*minimalism*)
 - ▶ ... **extensions** are formulated separately (*modularity*)
- ▶ ... without **dependencies** on other complex models (*standalone model*)
- ▶ ... without prescribing any domain-specific **vocabulary** (*separation from semantics*)
- ▶ ... not tied to any single **serialisation format** (*separation from syntax*)

STAM's **core model** is our answer to this.

- ▶ It is documented in a precise technical **specification** (*separation from implementation*)



Annotation

- (1) says **something** about **something**
- (2) any kind of **remark/classification/tagging** on a **text, a portion thereof, or on another annotation**

Figure 2: Annotations are central



Annotation

says **something** about **something**

any kind of **remark/classification/tagging** on a **text, a portion thereof, or on another annotation**

AnnotationData

Selector



Figure 3: AnnotationData and Selectors

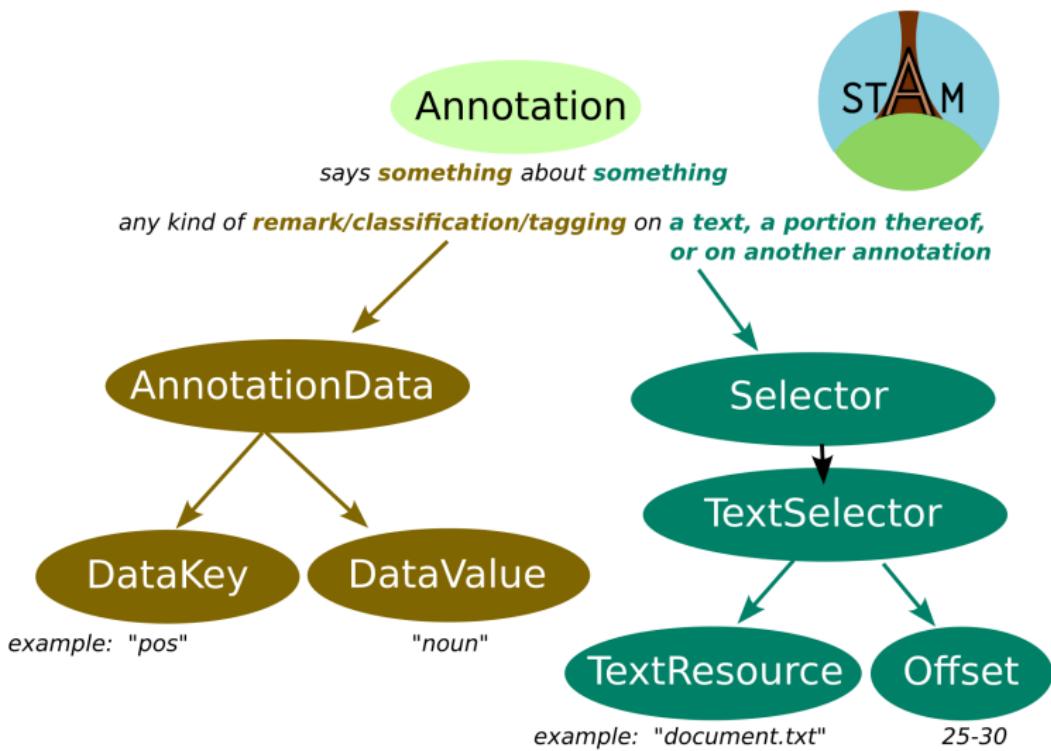


Figure 4: Keys, values and selector types

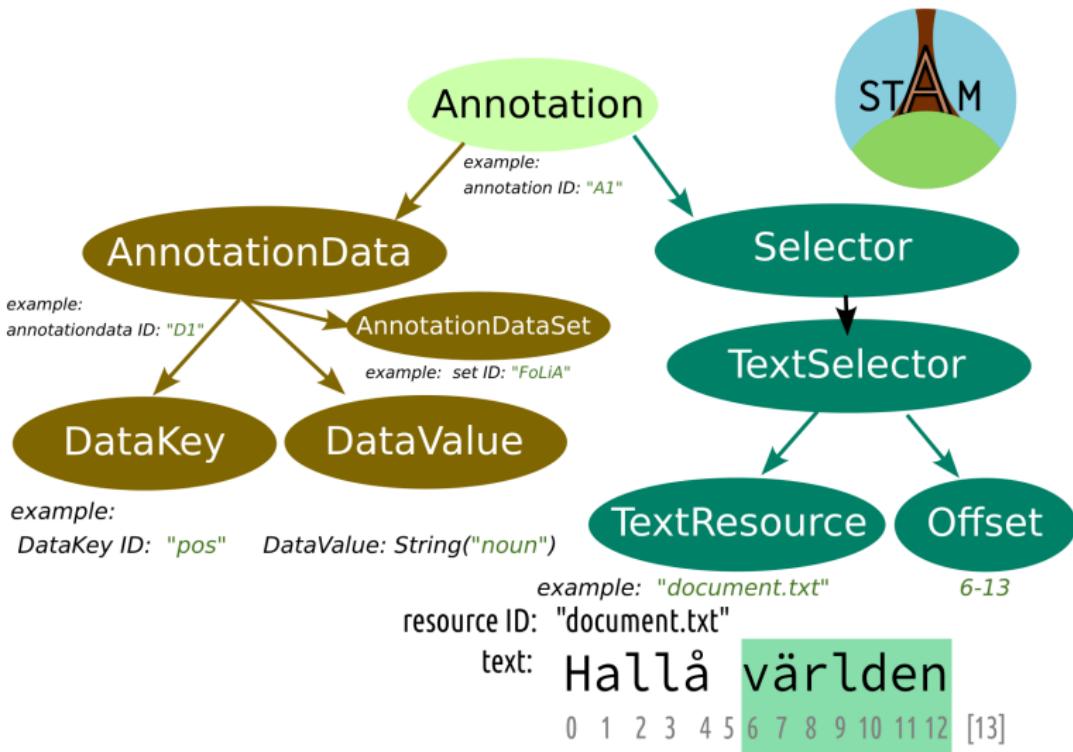


Figure 5: Example

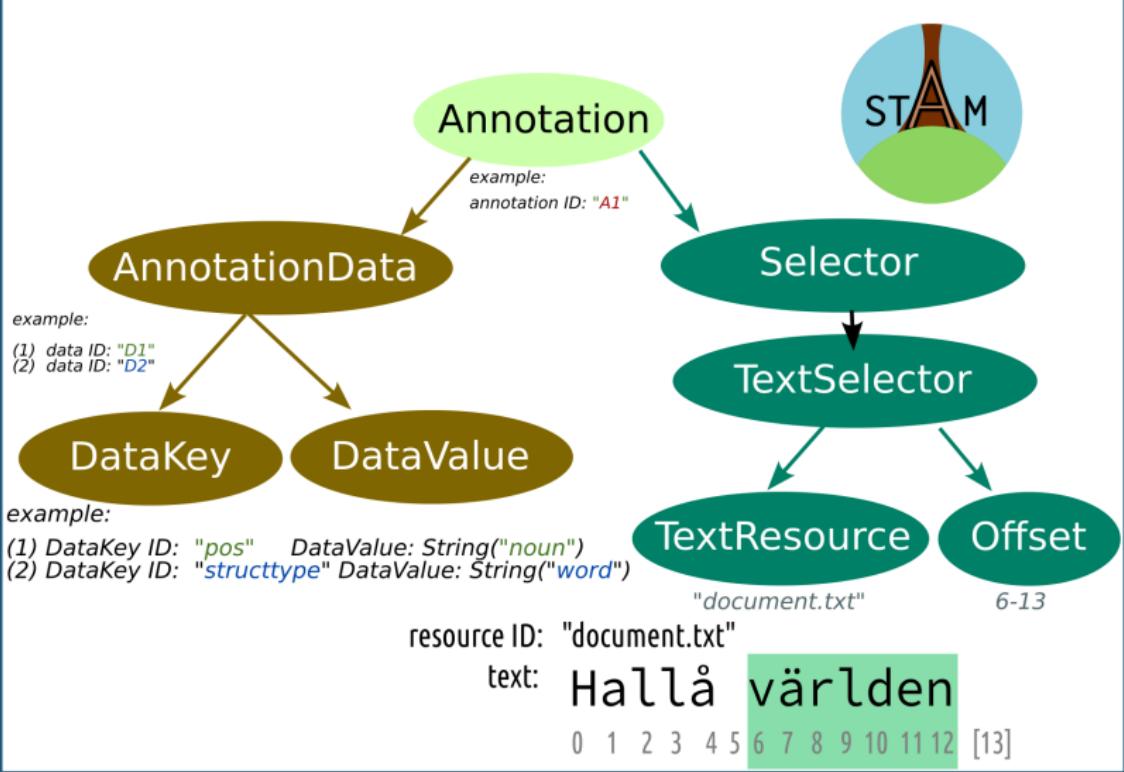


Figure 6: Example: Multiple AnnotationData

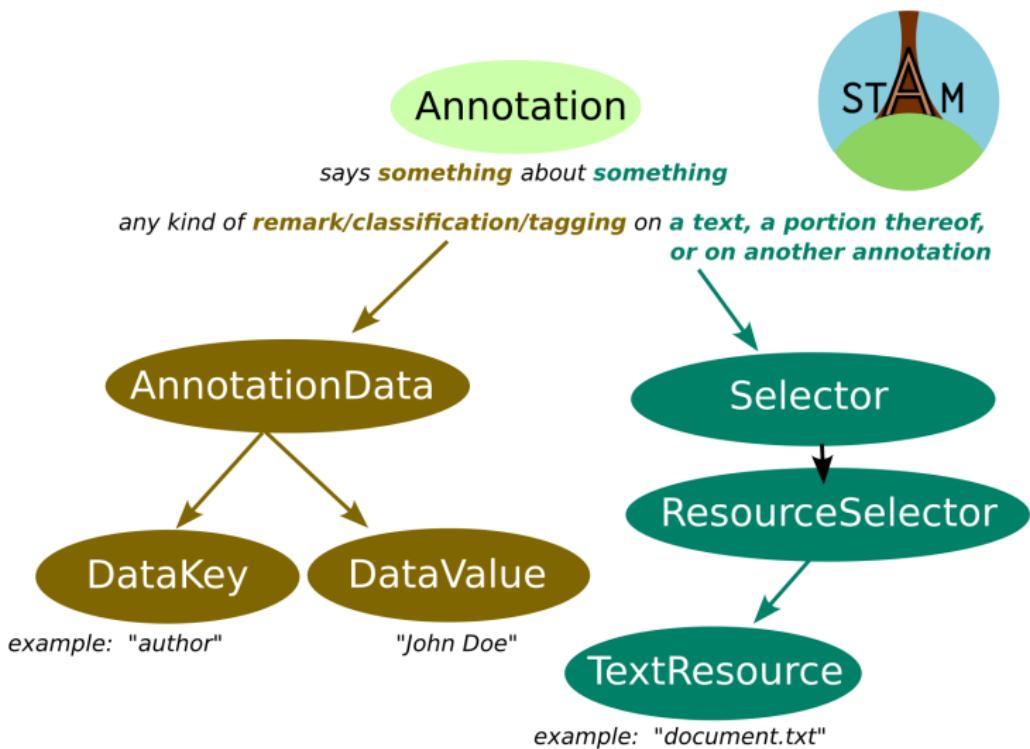


Figure 7: Example: ResourceSelector

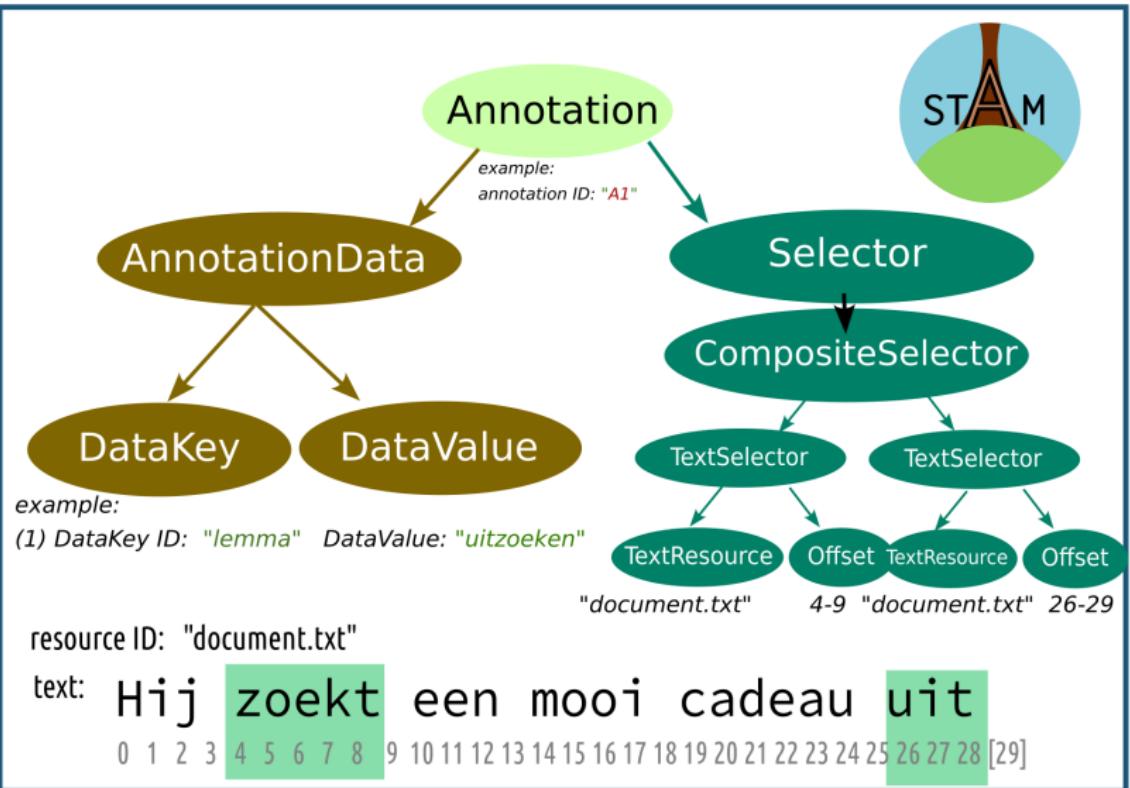


Figure 8: Example: CompositeSelector

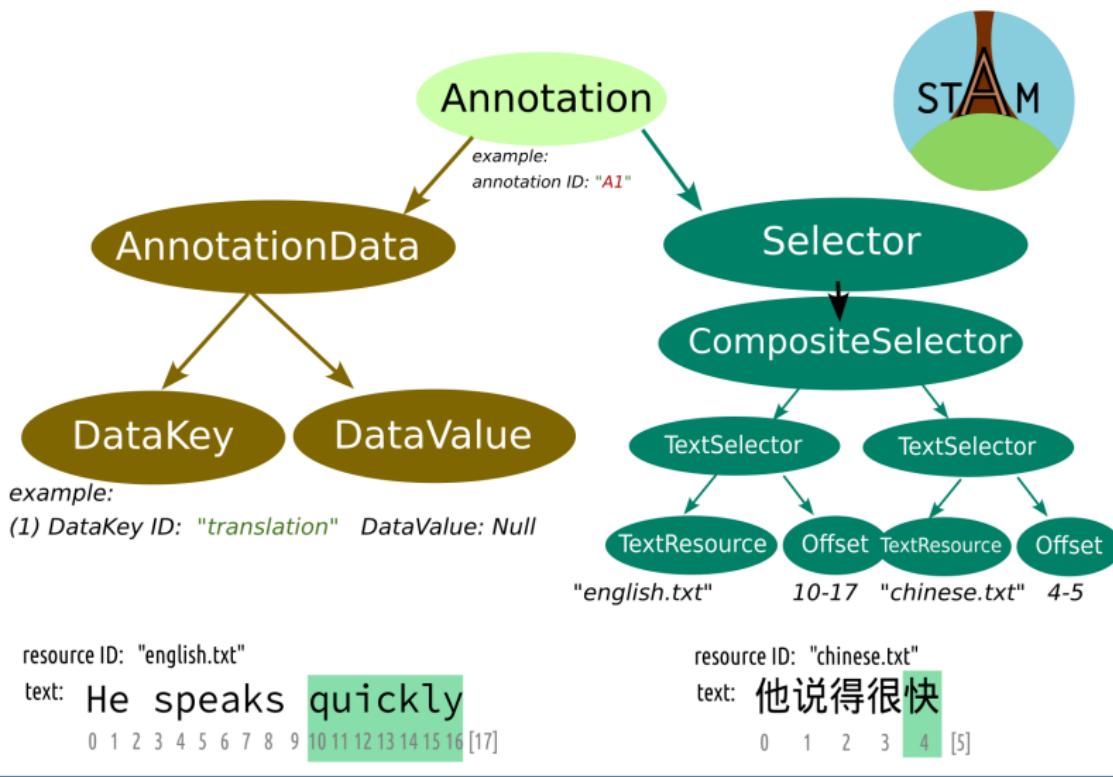


Figure 9: Example: CompositeSelector (2)

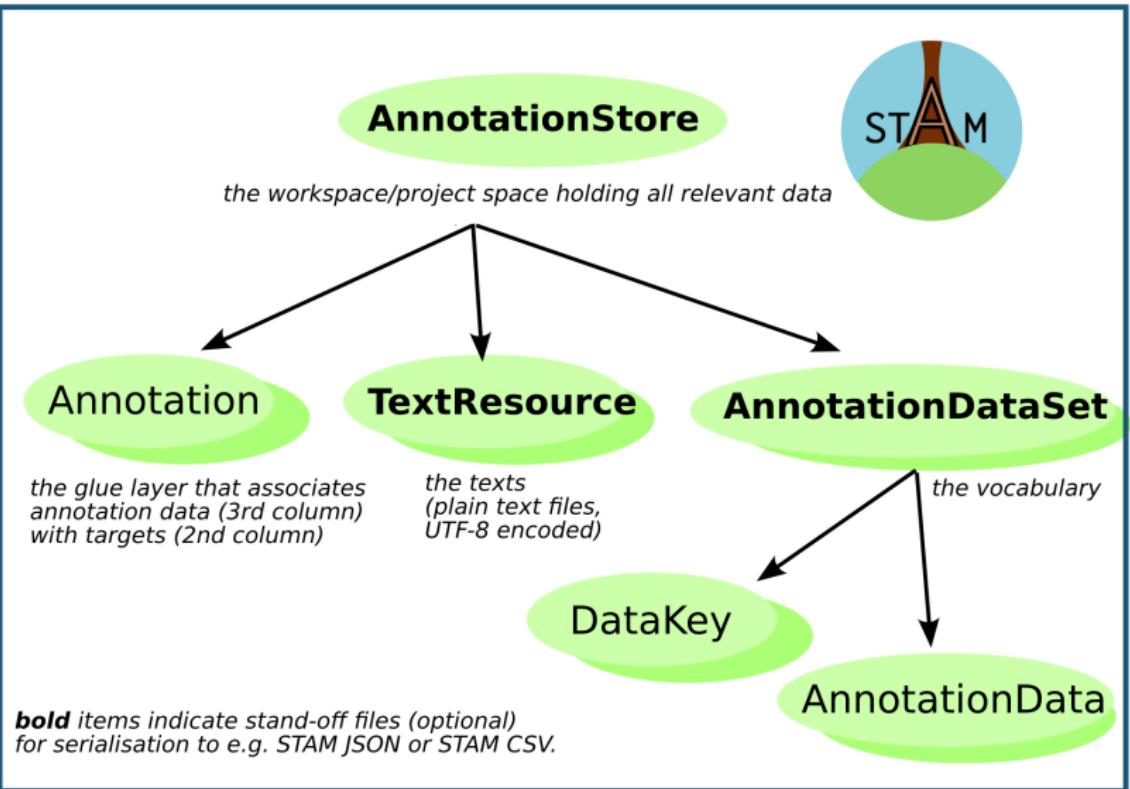


Figure 10: Collections

Data Model Characteristics

Technical:

- ▶ Directed Acyclic Graph
- ▶ Annotations should be regarded as immutable
- ▶ Optimised for efficient memory use
- ▶ Optimised for (reverse) search indices
- ▶ Almost everything is an annotation

STAM offers a framework to build upon, not the full solution:

- ▶ You have to choose **how to model** your data using STAM
- ▶ You have to choose what **vocabulary** to use for your data

Serialisation

	size	r/w speed	stand-off	readability	support
JSON	huge	normal	optional	good	canonical
CSV	medium	normal	always	good	extension
CBOR	small	fast	never	bad	implementation

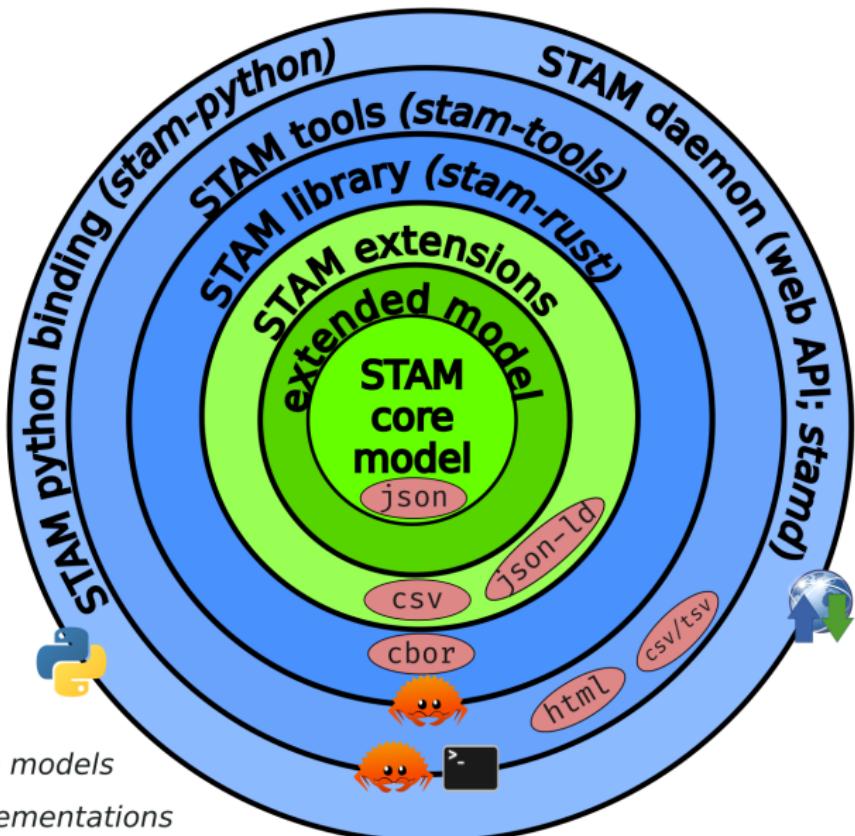
- ▶ *Stand-off.* Annotation Stores, Annotation Datasets, Text resources (plain text) can be stored in separate files.
- ▶ in the current STAM library (in-memory model), after parsing memory usage will be identical regardless of serialisation

```
[  
  {  
    "@type": "AnnotationStore", "@id": "example",  
    "resources": [ { "@type": "TextResource", "@id": "document.txt", "text": "Hallå världen" } ],  
    "annotationsets": [ { "@type": "AnnotationDataSet", "@id": "testset",  
      "keys": [ { "@type": "DataKey", "@id": "pos" } ],  
      "data": [ { "@type": "AnnotationData", "@id": "!D0", "key": "pos", "value": { "@type": "String", "value": "noun" } } ] } ],  
    "annotations": [  
      {  
        "@type": "Annotation", "@id": "!A0",  
        "target": {  
          "@type": "TextSelector",  
          "resource": "document.txt",  
          "offset": { "@type": "Offset", "begin": { "@type": "BeginAlignedCursor", "value": 6 },  
          | "end": { "@type": "BeginAlignedCursor", "value": 13 } }  
        },  
        "data": [ { "@type": "AnnotationData", "@id": "!D0", "set": "testset" } ]  
      }  
    ]  
  }]
```

Figure 11: STAM JSON

Tooling

- ▶ focus on **performance**: efficient usage of computing resources (memory, cpu)
- ▶ focus on **reusability**: implement once, and re-use, core logic needed for annotations on text
- ▶ **standalone**: runs on a wide variety of systems and independent of any wider service infrastructure: does NOT depend on any software services.
- ▶ **accessibility**: clean APIs/CLIs and thoroughly documented



Data models



Implementations



Serialisation formats



Interfaces

Implementations & Interfaces

Reference implementations and their interfaces:

- ▶ `stam-rust`: a high-performant library that implements the core model and all currently defined extensions
 - ▶ memory-based implementation, scalability constraints
 - ▶ reference implementation
- ▶ `stam-tools`: command line tools for working with STAM
 - ▶ unix principle: one tool (subcommand) for one thing
- ▶ `stam-python`: a Python-binding to `stam-rust` and `stam-tools`
 - ▶ good for technical data scientists
- ▶ `stamd`: STAM service offering a Web API

All target a technical audience and all are fully written in Rust
(compiles to native machine code)

Core functionality

The software library...

- ▶ ... exposes all concepts from the data model and holds data in memory efficiently
 - ▶ ... parses and serialises from/to disk
- ▶ ... allows editing data or creating it from scratch
- ▶ ... can retrieve texts for annotations, annotations for text, annotations for annotations
- ▶ ... can retrieve annotations based on *spatial* properties (*overlap, embedding, adjacency, proximity*)
- ▶ ... can retrieve annotations given annotation data (data indices)
- ▶ ... can search for text (regular expressions)
- ▶ ... can convert offset information between different reference frames
- ▶ ... can do validation

Functionality from extensions

- ▶ **STAM Vocab** – Express/validate against user-defined vocabularies
- ▶ **STAM Webannotations** – Allows modelling W3C Web Annotations using STAM (full export support)
- ▶ **STAM Textvalidation** – Adds an extra redundancy layer that helps protecting data integrity and aids readability of serialisations
- ▶ **STAM CSV** – Tabular serialisation format
- ▶ **STAM Query Language (STAMQL)** – An extensive query language that allows end-users to formulate and subsequently execute searches on a STAM model.
- ▶ **STAM Transpose** - Links identical textual parts across resources, allows *transposition* of annotations.
- ▶ **STAM Translate** - Links non-identical parts of texts, allows *translation* of annotations from one to another.

All extensions have their **own specifications** (*separation from implementation*)

Functionality in the tooling

Import & Export

- ▶ **CSV/TSV** with custom column layout (`stam import/stam export`):

Text	part-of-speech	part-of-sentence
I	pronoun	subject
see	verb	main verb
you	pronoun	object

- ▶ **W3C Web Annotations**; JSON-LD (`stam export, export only`)
- ▶ **XML** (`stam fromxml, import only`) – e.g. for *TEI XML*, *PageXML*, *FoLiA XML*
- ▶ **HTML** (`stam view, export only, visualisation`)



1. example.deep.p.1.s.1

De Russen kennen Nova Zembla sinds de 11e of 12e eeuw, toen
handelaars van Novgorod het eiland al aandeden.

2. example.deep.p.1.s.2

West-Europeanen ontdekeren de eilanden in de 16e eeuw tijdens de
zoektocht naar een noordoostelijke doorgang naar de Stille Oceaan.

Figure 13: STAM HTML visualisation (stam query -F html)

Querying

STAMQL is an *extensive* query language that exposes most functionality in the STAM library:

```
SELECT TEXT ?noun WHERE
    DATA "myset" "type" = "word";
    DATA "myset" "pos" = "noun"; {

SELECT TEXT ?adj WHERE
    RELATION ?noun SUCCEEDS;
    DATA "myset" "type" = "word";
    DATA "myset" "pos" = "adj";
}
```

Figure 14: A STAMQL query for adjective-noun word pairs

Querying and exporting often go hand-in-hand:

```
$ stam view --query 'SELECT RESOURCE ?res' \  
--query '@KEYVALUETAG SELECT ANNOTATION ?a WHERE  
RESOURCE ?res;  
DATA "testset" "pos" = "noun";' \  
example.stam.store.json
```

Hallå [världen|pos: noun]

Hallå världen pos: noun

Figure 15: A query and HTML visualisation via stam view

Searching and tagging text

- ▶ **Searching text** (`stam grep`) – Search for text (e.g. with regular expressions) and tag
- ▶ **Tagging** (`stam tag`) – ... and tag matches as annotations!

Alignment, Transposition & Translation

- ▶ **Alignment** – (stam align) Align two similar texts (Smith-Waterman / Needleman-Wunsch)
 - ▶ produces transpositions or translations
- ▶ **Normalisation** – (stam translatetext) Change a text based on simple rewrite rules
 - ▶ produces translations
- ▶ **Translation & Transposition** – Copy annotations from one text to another over a translation/transposition
 - ▶ stam translate & stam transpose

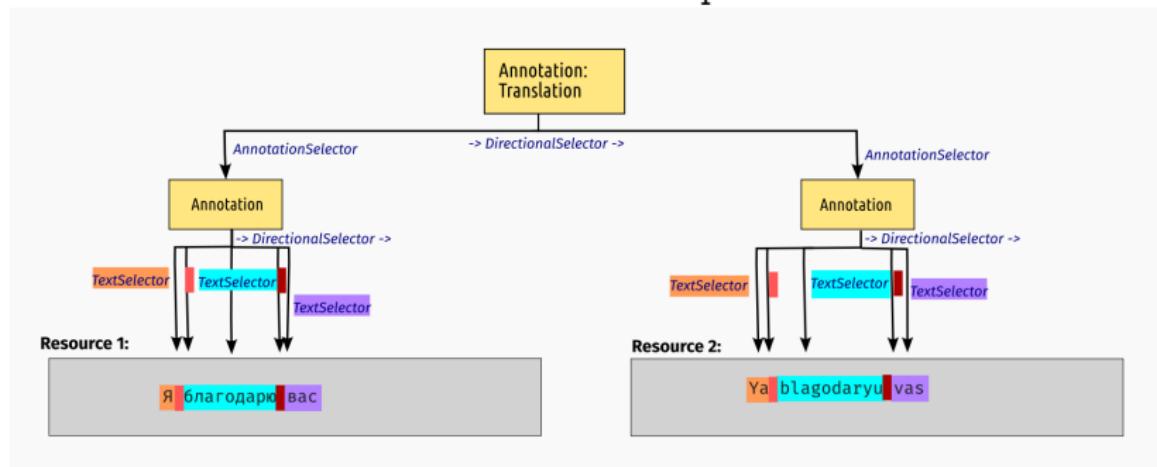


Figure 16: A translation annotation

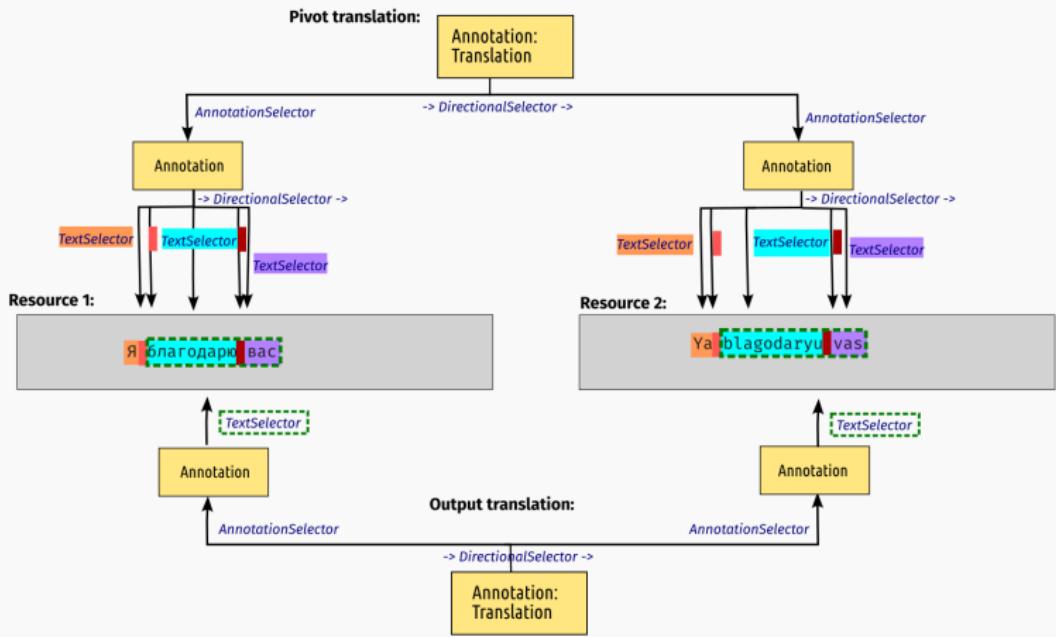


Figure 17: Copying an annotation over a pivot translation

Other functionality

- ▶ **Introspection** – View the raw data in the model and its memory consumption
- ▶ **Merging and splitting annotation stores**
- ▶ **Validation**

Conclusion

STAM ...

- ▶ ... offers a **solid generic foundation** upon which you can build higher-level applications for working with stand-off annotation on text
- ▶ ... is **modular**; pick and use the parts you need for your task
- ▶ ... is **flexible**; you decide how to model your data and bring **your own vocabulary**
- ▶ ... takes care of the boring **computational groundwork** so you don't have to
- ▶ ... provides different **interfaces** for different technical **audiences**
- ▶ ... has focus on **performance** and implementations are written in Rust
- ▶ ... is **free open source software** (GNU General Public License v3) and open to contributors.
- ▶ ... > 40k lines Rust code, > 1500 development hours, since january 2023 (funded in part by CLARIAH)

Fin

Further information: Project website –

<https://annotation.github.io/stam>

- ▶ STAM Specification & Extension specifications
- ▶ Library API References (Rust & Python)
- ▶ Python tutorial “Standoff Text Annotation for Pythonistas”
(Jupyter Notebook)
- ▶ Screencast video for stam-tools
- ▶ Poster
- ▶ These slides

